

## ✓ SAT 5114

### AI in HEALTHCARE PROJECT

## ✓ Install and load libraries

```
!pip install -qqq "unsloth[colab-new] @ git+https://github.com/unslothai/unsloth.git" --progress-bar off
!pip install -qqq xformers trl peft accelerate bitsandbytes triton --progress-bar off
!pip install -qqq unsloth transformers accelerate datasets peft bitsandbytes wandb evaluate bert-score
```

```

Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
Building wheel for unsloth (pyproject.toml) ... done
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source
gcsfs 2025.3.2 requires fsspec==2025.3.2, but you have fsspec 2024.12.0 which is incompatible.
tensorflow-metadata 1.17.1 requires protobuf<6.0.0,>=4.25.2; python_version >= "3.11", but you have protobuf 3.20.3 which is incompatibl
grpcio-status 1.71.0 requires protobuf<6.0dev,>=5.26.1, but you have protobuf 3.20.3 which is incompatible.
ydf 0.11.0 requires protobuf<6.0.0,>=5.29.1, but you have protobuf 3.20.3 which is incompatible.
===== 84.0/84.0 kB 2.5 MB/s eta 0:00:00
===== 61.1/61.1 kB 6.1 MB/s eta 0:00:00

```

## ✓ Import Libraries

```

# Import libraries
import torch
import unsloth
import numpy as np
from tqdm import tqdm
from datasets import load_dataset
from transformers import (
    AutoModelForCausalLM,
    AutoTokenizer,
    TrainingArguments,
    Trainer,
    EarlyStoppingCallback,
    TextStreamer,
    LogitsProcessor,
    LogitsProcessorList,
    BitsAndBytesConfig,
)
from peft import LoraConfig, get_peft_model
from sklearn.metrics import f1_score
import evaluate

```

```

🚀 Unsloth: Will patch your computer to enable 2x faster free finetuning.
Unsloth: Failed to patch Gemma3ForConditionalGeneration.
🚀 Unsloth Zoo will now patch everything to make training faster!

```

## ✓ Install and define evaluation

```

!pip install rouge_score
bleu_metric = evaluate.load("bleu")
rouge_metric = evaluate.load("rouge")
bertscore_metric = evaluate.load("bertscore")

```

```

Collecting rouge_score
  Downloading rouge_score-0.1.2.tar.gz (17 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: absl-py in /usr/local/lib/python3.11/dist-packages (from rouge_score) (1.4.0)
Requirement already satisfied: nltk in /usr/local/lib/python3.11/dist-packages (from rouge_score) (3.9.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from rouge_score) (2.0.2)
Requirement already satisfied: six>=1.14.0 in /usr/local/lib/python3.11/dist-packages (from rouge_score) (1.17.0)
Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk->rouge_score) (8.1.8)
Requirement already satisfied: joblib in /usr/local/lib/python3.11/dist-packages (from nltk->rouge_score) (1.4.2)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk->rouge_score) (2024.11.6)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk->rouge_score) (4.67.1)
Building wheels for collected packages: rouge_score
  Building wheel for rouge_score (setup.py) ... done
  Created wheel for rouge_score: filename=rouge_score-0.1.2-py3-none-any.whl size=24934 sha256=8b15b77f151f7057b1fba1daa2b39b6e5a9e9489e
  Stored in directory: /root/.cache/pip/wheels/1e/19/43/8a442dc83660ca25e163e1bd1f89919284ab0d0c1475475148
Successfully built rouge_score
Installing collected packages: rouge_score
Successfully installed rouge_score-0.1.2

Downloading builder script: 100% 5.94k/5.94k [00:00<00:00, 649kB/s]

Downloading extra modules: 4.07k/? [00:00<00:00, 337kB/s]

Downloading extra modules: 100% 3.34k/3.34k [00:00<00:00, 298kB/s]

Downloading builder script: 100% 6.27k/6.27k [00:00<00:00, 533kB/s]

Downloading builder script: 100% 7.95k/7.95k [00:00<00:00, 818kB/s]

```

## ✓ Load and prepare the dataset

```

dataset = load_dataset("Shekswess/gemma_medquad_instruct_dataset", split="train[:400]")
dataset = dataset.train_test_split(test_size=0.1)

```

```

#lavita/ChatDoctor-HealthCareMagic-100k
#Shekswess/gemma_medquad_instruct_dataset

```

```

README.md: 100% 1.64k/1.64k [00:00<00:00, 197kB/s]

train-00000-of-00001.parquet: 100% 17.9M/17.9M [00:00<00:00, 55.3MB/s]

Generating train split: 100% 16359/16359 [00:00<00:00, 170279.88 examples/s]

```

## ✓ Format the dataset

```

# Formatting function
def format_instruction(example):
    return {
        "prompt": f"### Instruction:\n{example['instruction']}\n\n### Input:\n{example['input']}\n\n### Response:\n",
        "response": example["output"]
    }

```

```
dataset = dataset.map(format_instruction)
```

```

Map: 100% 360/360 [00:00<00:00, 6465.04 examples/s]

Map: 100% 40/40 [00:00<00:00, 2145.51 examples/s]

```

## ✓ Model setup and quantization

```

# Model setup with 4-bit quantization
model_name = "unsloth/llama-3-8b"
quantization_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_quant_type="nf4",
    bnb_4bit_compute_dtype=torch.float16,
    bnb_4bit_use_double_quant=True,
)

```

```

tokenizer = AutoTokenizer.from_pretrained(model_name)
tokenizer.pad_token = tokenizer.eos_token

```

```

model = AutoModelForCausalLM.from_pretrained(
    model_name,
    quantization_config=quantization_config,
    device_map="auto",
    use_cache=False,
)

```



Loading checkpoint shards: 100%

4/4 [00:18&lt;00:00, 3.82s/it]

## ✓ LoRA configuration and PEFT

```

# Improved LoRA configuration
peft_config = LoraConfig(
    r=4,
    lora_alpha=8,
    target_modules=["q_proj", "v_proj", "k_proj", "o_proj"],
    lora_dropout=0.05,
    bias="lora_only",
    task_type="CAUSAL_LM",
    inference_mode=False,
)
model = get_peft_model(model, peft_config)
model.enable_input_require_grads()

```

## ✓ Data Collation and Tokenization

```

# Add data collator
from transformers import DataCollatorForSeq2Seq

data_collator = DataCollatorForSeq2Seq(
    tokenizer,
    pad_to_multiple_of=8,
    padding=True,
    return_tensors="pt",
)

# Update tokenize function
def tokenize_function(examples):
    texts = [p + r for p, r in zip(examples["prompt"], examples["response"])]
    tokenized = tokenizer(
        texts,
        max_length=512,
        truncation=True,
        padding="max_length",
        add_special_tokens=False
    )
    return {
        "input_ids": tokenized["input_ids"],
        "attention_mask": tokenized["attention_mask"],
        "labels": tokenized["input_ids"].copy()
    }

tokenized_dataset = dataset.map(tokenize_function, batched=True)

```



Map: 100%

360/360 [00:00&lt;00:00, 2158.53 examples/s]

Map: 100%

40/40 [00:00&lt;00:00, 1231.32 examples/s]

## ✓ Metrics calculation

```

# Metrics calculation (updated)
def compute_metrics(eval_pred):
    preds, labels = eval_pred
    labels = np.where(labels != -100, labels, tokenizer.pad_token_id)

    # Convert logits to token IDs (shape: [batch_size, seq_length])
    preds = np.argmax(preds, axis=-1) # Add this line

```

```

pred_texts = tokenizer.batch_decode(preds, skip_special_tokens=True)
label_texts = tokenizer.batch_decode(labels, skip_special_tokens=True)

results = {}

# BLEU
results["bleu"] = bleu_metric.compute(
    predictions=pred_texts,
    references=[[text] for text in label_texts]
)["bleu"]

# ROUGE
results.update(rouge_metric.compute(
    predictions=pred_texts,
    references=label_texts
))

# BERTScore
bert_results = bertscore_metric.compute(
    predictions=pred_texts,
    references=label_texts,
    lang="en"
)
results["bert_score"] = np.mean(bert_results["f1"])

return results

```

## ✓ Setup Training arguments and trainer


```

# Modified TrainingArguments with evaluation strategy
training_args = TrainingArguments(
    output_dir="./llama3_healthcare",
    per_device_train_batch_size=2,
    per_device_eval_batch_size=4,
    gradient_accumulation_steps=1,
    save_strategy='epoch',          # Changed to "epoch" to enable saving
    logging_strategy="no",          # Disable logging
    learning_rate=1e-5,
    weight_decay=1,
    num_train_epochs=4,
    lr_scheduler_type="cosine",
    warmup_ratio=0.1,
    logging_steps=50,
    load_best_model_at_end=True,
    fp16=True,
    gradient_checkpointing=True,
    report_to="none",               # Disabled all reporting
    remove_unused_columns=False,
    # Add evaluation strategy for EarlyStoppingCallback
    eval_strategy = "epoch" # or "steps" with logging_steps defined
)

# Trainer remains the same
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_dataset["train"],
    eval_dataset=tokenized_dataset["test"],
    compute_metrics=compute_metrics,
    callbacks=[EarlyStoppingCallback(early_stopping_patience=3)]
)

# Training will now only log to console
trainer.train()

```



[720/720 06:31, Epoch 4/4]

Epoch	Training Loss	Validation Loss	Bleu	Rouge1	Rouge2	RougeL	RougeLsum	Bert Score
1	No log	0.809389	0.377964	0.685381	0.374715	0.575169	0.605812	0.883580
2	No log	0.665850	0.452822	0.738861	0.455310	0.651004	0.665704	0.905270
3	No log	0.642137	0.459082	0.738626	0.456576	0.652674	0.665940	0.905385
4	No log	0.639166	0.459274	0.739225	0.457579	0.652887	0.666850	0.905519

tokenizer\_config.json: 100%

25.0/25.0 [00:00<00:00, 2.99kB/s]

config.json: 100%

482/482 [00:00<00:00, 61.0kB/s]

vocab.json: 100%

899k/899k [00:00<00:00, 16.8MB/s]

merges.txt: 100%

456k/456k [00:00<00:00, 12.4MB/s]

tokenizer.json: 100%

1.36M/1.36M [00:00<00:00, 19.0MB/s]

Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download. For better performance, you should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.  
WARNING:huggingface\_hub.file\_download:Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download.

model.safetensors: 100%

1.42G/1.42G [00:02<00:00, 1.34GB/s]

Some weights of RobertaModel were not initialized from the model checkpoint at roberta-large and are newly initialized: ['pooler.dense.tanh\_proj', 'pooler.dense.linear']. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.  
TrainOutput(global\_step=720, training\_loss=1.4614861382378472, metrics={'train\_runtime': 392.4554, 'train\_samples\_per\_second': 3.669, 'train\_steps\_per\_second': 1.835, 'total\_flos': 3.321446050824192e+16, 'train\_loss': 1.4614861382378472, 'epoch': 4.0})

Model saving

```
# Save and generation code remains unchanged
model.save_pretrained("./llama3_healthcare_finetuned")
tokenizer.save_pretrained("./llama3_healthcare_finetuned")

('./llama3_healthcare_finetuned/tokenizer_config.json',
 './llama3_healthcare_finetuned/special_tokens_map.json',
 './llama3_healthcare_finetuned/tokenizer.json')
```

Text generation

```
from transformers import pipeline # Import the pipeline function
generator = pipeline(
    "text-generation",
    model=model,
    tokenizer=tokenizer,

    max_new_tokens=256,
    do_sample=True,
    temperature=0.7,
    top_p=0.9,
    repetition_penalty=1.2,
)

test_question = "What is the management for hypertension?"
result = generator(test_question, num_return_sequences=1)
print("\nGenerated Response:")
print(result[0]['generated_text'])

Device set to use cuda:0
The model 'PefitModelForCausalLM' is not supported for text-generation. Supported models are ['AriaTextForCausalLM', 'BambaForCausalLM',
Generated Response:
What is the management for hypertension? How to control it?
Hypertension can be controlled by a combination of lifestyle changes and medications. It's important that you make these modifications to
Lifestyle Changes - The first line treatment should always include dietary modification (i.e., reduction in salt intake), physical activity, and
Medications - If blood pressure remains uncontrolled after three months despite following all advice above then doctors usually recommend
A variety of different medicines exist which treat high BP including diuretics (water pills); beta-blockers like Atenolol or Metoprolol;
```

Start coding or generate with AI.

- ✓ Add Graphical User Interface for input and output texts
- ✓ Install GUI library gradio

```
!pip install gradio
```

 [Show hidden output](#)

- ✓ Define prediction Function

```
import gradio as gr

def predict(input_text):
    result = generator(input_text, num_return_sequences=1)
    return result[0]['generated_text']
```

- ✓ Create the user interface

```
iface = gr.Interface(
    fn=predict,
    inputs=gr.Textbox(lines=2, placeholder="Enter your question here..."),
    outputs="text",
    title="Medical Question Answering",
    description="Ask questions about medical topics and get answers from our AI model."
)
```

```
iface.launch()
```

 Rerunning server... use `close()` to stop if you need to change `launch()` parameters.

----

It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sharing must be enabled. Automatically

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

\* Running on public URL: <https://ecef7e90590802.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working dir

## Medical Question Answering

Ask questions about medical topics and get answers from our AI model.

input\_text

What is the management for hypertension

Clear

Submit

output

What is the management for hypertension?

Hypertension or high blood pressure can be controlled and managed by following a healthy lifestyle. This includes:

Healthy diet – A low salt, rich in fruits vegetables and nuts diet helps to control your blood pressure levels.

Regular exercise – Exercising regularly keeps you fit and controls your weight which has been linked with increased risk of developing heart disease.

Avoid smoking- Smoking causes narrowing of arteries leading to development of cardiovascular diseases like coronary artery disease (CAD) or stroke

Maintain normal body weight – Weight gain increases your chances of having diabetes along with other problems such as joint pain etc., so it's important not only that one must lose extra pounds but also maintain them at an ideal level through regular physical activity combined with balanced eating habits